

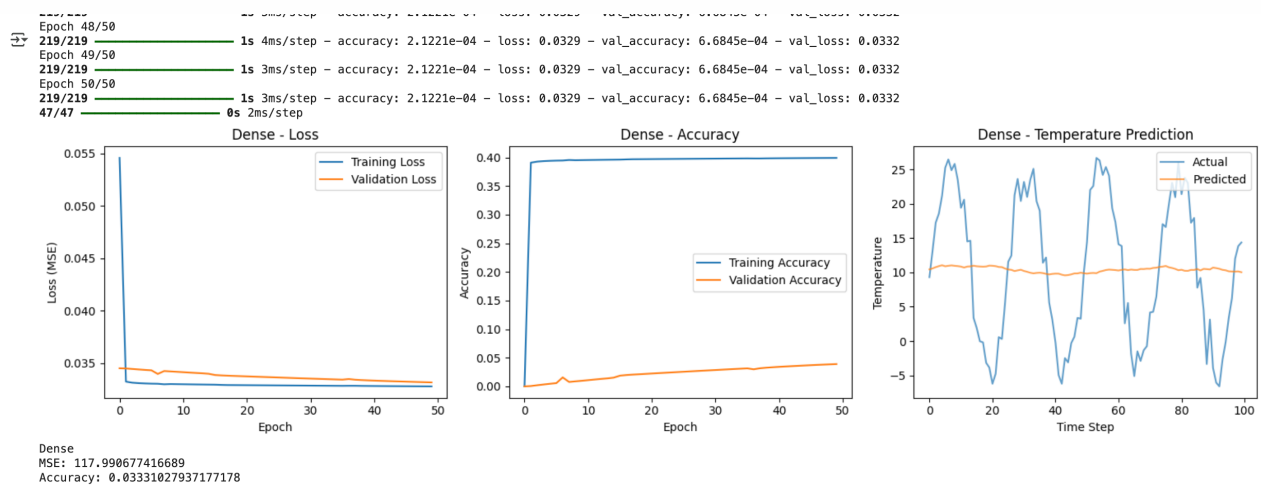
# Recurrent Neural Network Report

Yuhuan Huang, 2025

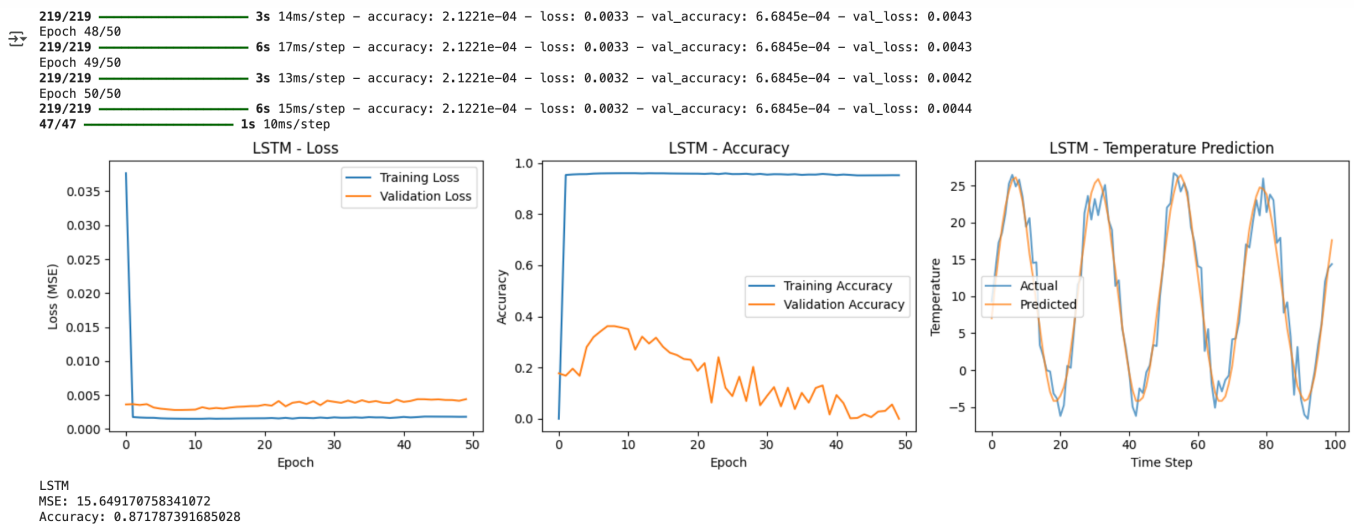
## Part1: LSTM on synthetic datasets

Console output:

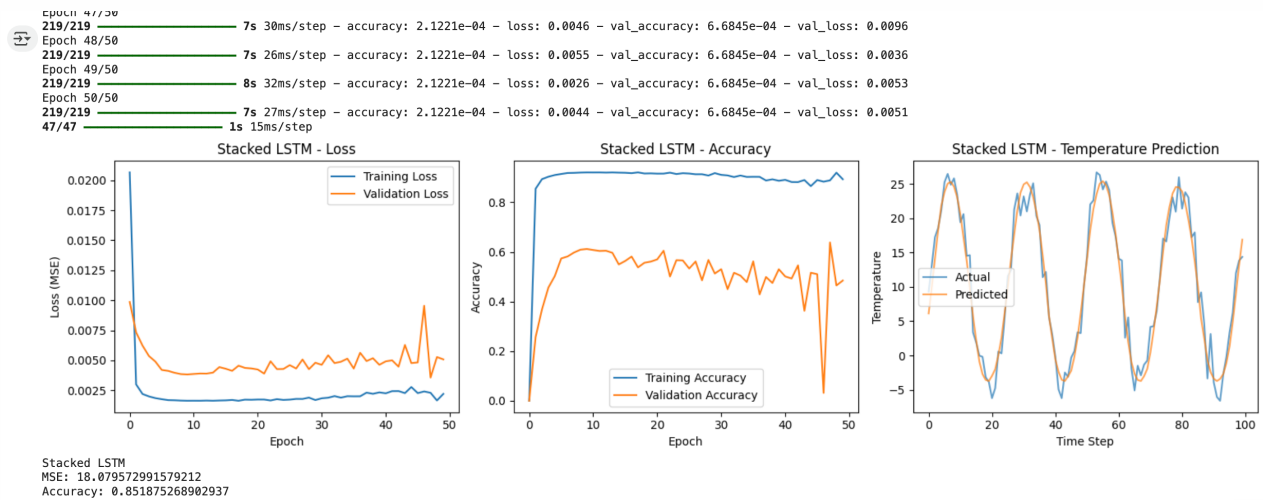
### 1. Dense:



### 2. LSTM:



### 3. Stacked-LSTM (two-layer LSTM):



```
Stacked LSTM
MSE: 18.079572991579212
Accuracy: 0.851875268902937
Dense MSE: 117.990677416689
Dense Accuracy: 0.03331027937177178
LSTM MSE: 15.649170758341072
LSTM Accuracy: 0.871787391685028
Stacked LSTM MSE: 18.079572991579212
Stacked LSTM Accuracy: 0.851875268902937
*****
File "__main__", line 3, in __main__
Failed example:
  np.isclose(results['Dense']['mse'], 117.33981323767615, atol=1e-02)
Expected:
  np.True_
Got:
  np.False_
*****
1 items had failures:
  1 of 6 in __main__
***Test Failed*** 1 failures.
TestResults(failed=1, attempted=6)
```

I am sorry that the MSE test failed, and I don't know how to fix it. From all the above models, I can see the performance: LSTM > Stacked LSTM > Dense

## Reflection:

1. The Dense neural network performed poorly compared to both LSTM models. Why might a Dense network struggle with time series prediction, even though it has access to the same 24-hour window of temperature data?

**Ans:** Dense is not a time-series model, which means that it would treat the 24-hour data as each independent points, but not time series, so it would lose information given by the periodicity of a time series.

2. The Stacked LSTM (two LSTM layers) performed slightly worse than the single LSTM, despite being a more complex model. What are some potential reasons for this, and what does it tell us about choosing model architectures?

**Ans:** Two-layer LSTM may make the model too complex, less efficient, or cause over-fitting problem. So when we are choosing models, it is not always better to choose the more complex model.

3. Looking at the temperature predictions plot, you'll notice that all models (even the poor-performing Dense network) capture some basic pattern in the data. What might each model be learning about temperature patterns, and how does this relate to their architectures?

**Ans:** The LSTM and Stacked LSTM can capture the periodicity of the temperature data, while the Dense model can only capture the general trend (look like "average") of the temperature data. This is because Dense data is not a time-series model, but LSTM models are time-series models.

## Part2: NLP Architecture Comparisons

### 1. The Sarcasm dataset:

```
... Sarcasm dataset configurations:
Learning rates: {'Dense Network': 0.0001, 'LSTM Network': 1e-05, 'Bidirectional LSTM Network': 1e-05, 'Double Bidirectional LSTM Network': 1e-05}
Embedding dimensions: {'Dense Network': 16, 'LSTM Network': 64, 'Bidirectional LSTM Network': 64, 'Double Bidirectional LSTM Network': 64}
Epochs: {'Dense Network': 30, 'LSTM Network': 20, 'Bidirectional LSTM Network': 15, 'Double Bidirectional LSTM Network': 15}

Running comparison on Sarcasm dataset with optimized hyperparameters:
Dense Network: embedding_dim=16, epochs=30, lr=0.0001
LSTM Network: embedding_dim=64, epochs=20, lr=1e-05
Bidirectional LSTM Network: embedding_dim=64, epochs=15, lr=1e-05
Double Bidirectional LSTM Network: embedding_dim=64, epochs=15, lr=1e-05

Initial configuration:
Learning rates: {'Dense Network': 0.0001, 'LSTM Network': 1e-05, 'Bidirectional LSTM Network': 1e-05, 'Double Bidirectional LSTM Network': 1e-05}
Embedding dimensions: {'Dense Network': 16, 'LSTM Network': 64, 'Bidirectional LSTM Network': 64, 'Double Bidirectional LSTM Network': 64}
Epochs config: {'Dense Network': 30, 'LSTM Network': 20, 'Bidirectional LSTM Network': 15, 'Double Bidirectional LSTM Network': 15}

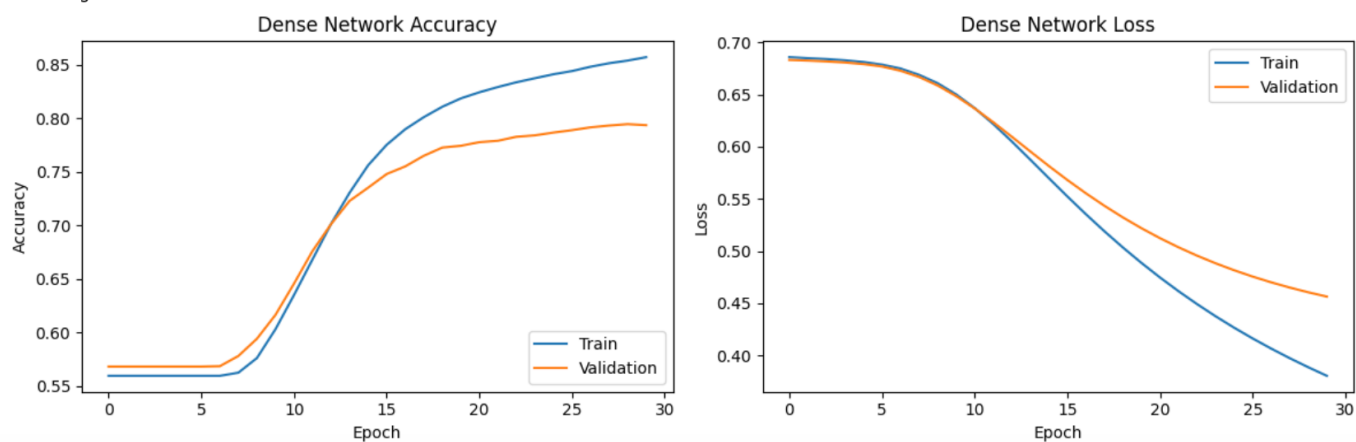
Final configuration after defaults:
Learning rates: {'Dense Network': 0.0001, 'LSTM Network': 1e-05, 'Bidirectional LSTM Network': 1e-05, 'Double Bidirectional LSTM Network': 1e-05}
Embedding dimensions: {'Dense Network': 16, 'LSTM Network': 64, 'Bidirectional LSTM Network': 64, 'Double Bidirectional LSTM Network': 64}
Epochs config: {'Dense Network': 30, 'LSTM Network': 20, 'Bidirectional LSTM Network': 15, 'Double Bidirectional LSTM Network': 15}

=====
Running Experiment on sarcasm Dataset
=====
Loading Sarcasm dataset...
Preprocessing data...
Training set size: 21367
Validation set size: 5342
```

Dense:

- Epoch 28/30  
668/668 ————— 4s 4ms/step - accuracy: 0.8537 - loss: 0.3998 - val\_accuracy: 0.7931 - val\_loss: 0.4650
- Epoch 29/30  
668/668 ————— 3s 4ms/step - accuracy: 0.8566 - loss: 0.3909 - val\_accuracy: 0.7945 - val\_loss: 0.4605
- Epoch 30/30  
668/668 ————— 5s 5ms/step - accuracy: 0.8600 - loss: 0.3824 - val\_accuracy: 0.7935 - val\_loss: 0.4564

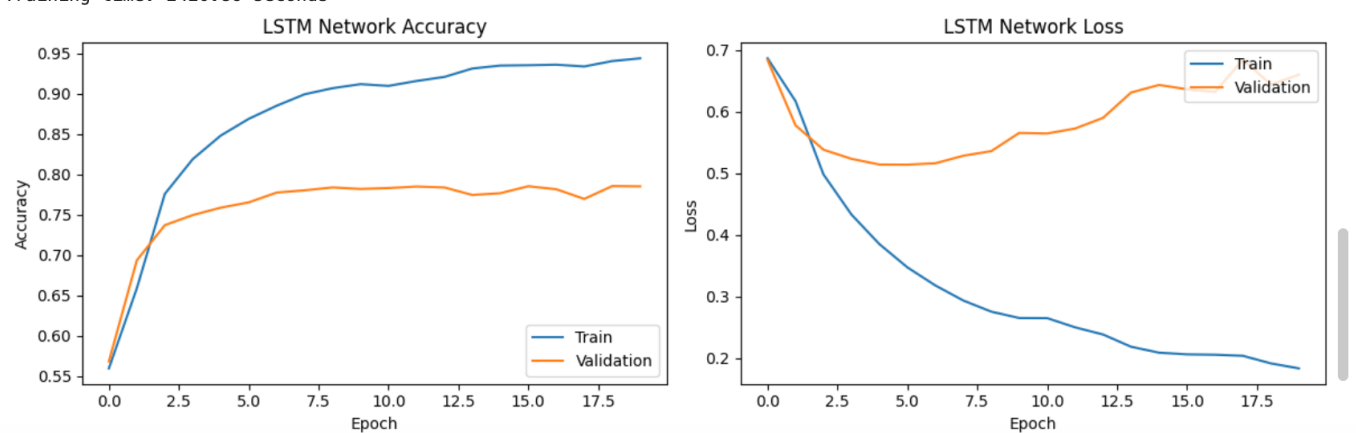
Dense Network Validation Accuracy: 0.7935  
Training time: 122.17 seconds



## LSTM:

- Epoch 18/20  
668/668 ————— 138s 163ms/step - accuracy: 0.9287 - loss: 0.2124 - val\_accuracy: 0.7696 - val\_loss: 0.6844
- Epoch 19/20  
668/668 ————— 106s 158ms/step - accuracy: 0.9341 - loss: 0.2013 - val\_accuracy: 0.7855 - val\_loss: 0.6439
- Epoch 20/20  
668/668 ————— 145s 162ms/step - accuracy: 0.9402 - loss: 0.1884 - val\_accuracy: 0.7851 - val\_loss: 0.6601

LSTM Network Validation Accuracy: 0.7851  
Training time: 2420.86 seconds



## Bidirectional LSTM

Epoch 13/15

**668/668** ————— **83s** 114ms/step - accuracy: 0.8380 - loss: 0.4059 - val\_accuracy: 0.7819 - val\_loss: 0.4717

Epoch 14/15

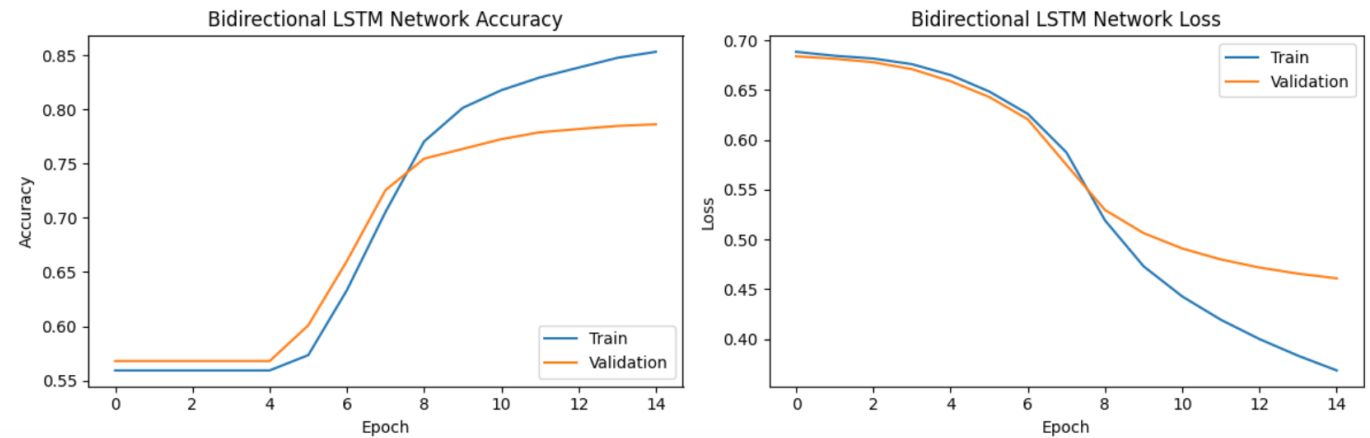
**668/668** ————— **76s** 114ms/step - accuracy: 0.8471 - loss: 0.3885 - val\_accuracy: 0.7847 - val\_loss: 0.4655

Epoch 15/15

**668/668** ————— **76s** 114ms/step - accuracy: 0.8518 - loss: 0.3732 - val\_accuracy: 0.7862 - val\_loss: 0.4608

Bidirectional LSTM Network Validation Accuracy: 0.7862

Training time: 1172.97 seconds



## Double Bidirectional LSTM

**668/668** ————— **160s** 240ms/step - accuracy: 0.8757 - loss: 0.3205 - val\_accuracy: 0.7903 - val\_loss: 0.4778

Epoch 10/15

**668/668** ————— **161s** 241ms/step - accuracy: 0.8863 - loss: 0.2984 - val\_accuracy: 0.7917 - val\_loss: 0.4853

Epoch 11/15

**668/668** ————— **160s** 239ms/step - accuracy: 0.8957 - loss: 0.2802 - val\_accuracy: 0.7915 - val\_loss: 0.4946

Epoch 12/15

**668/668** ————— **159s** 238ms/step - accuracy: 0.9046 - loss: 0.2644 - val\_accuracy: 0.7892 - val\_loss: 0.5056

Epoch 13/15

**668/668** ————— **160s** 240ms/step - accuracy: 0.9112 - loss: 0.2502 - val\_accuracy: 0.7900 - val\_loss: 0.5181

Epoch 14/15

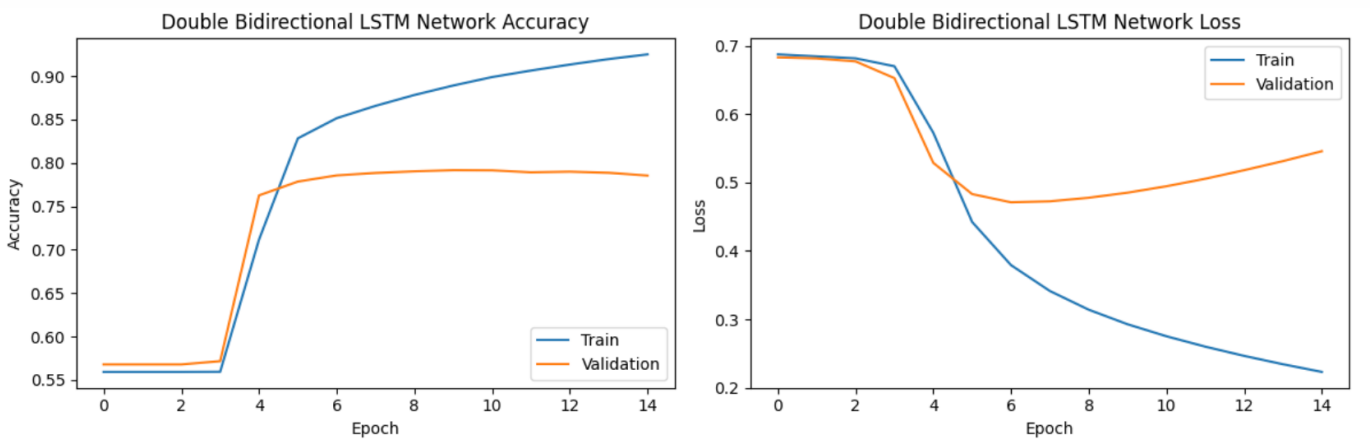
**668/668** ————— **161s** 241ms/step - accuracy: 0.9177 - loss: 0.2372 - val\_accuracy: 0.7887 - val\_loss: 0.5314

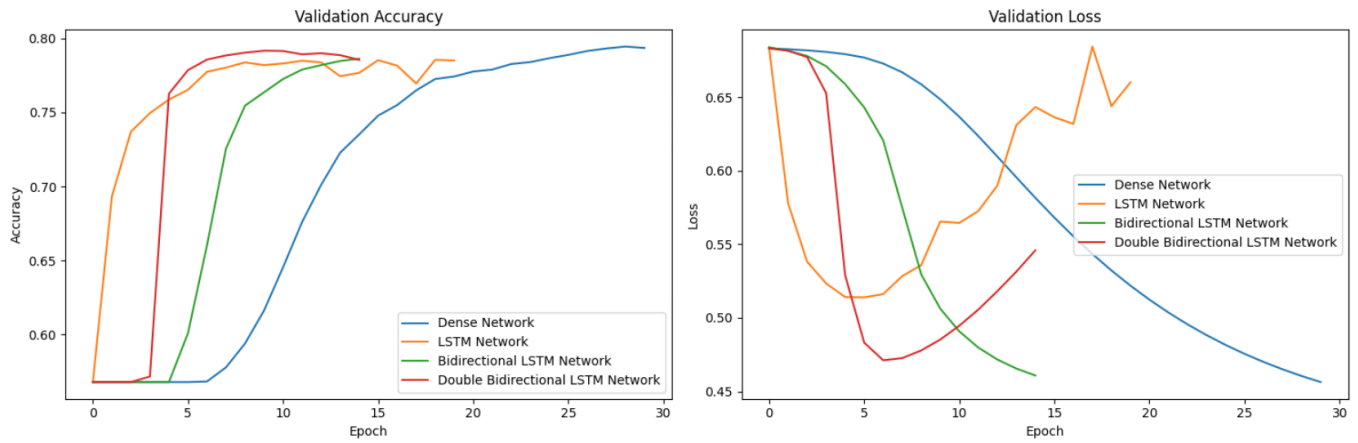
Epoch 15/15

**668/668** ————— **160s** 240ms/step - accuracy: 0.9236 - loss: 0.2255 - val\_accuracy: 0.7855 - val\_loss: 0.5459

Double Bidirectional LSTM Network Validation Accuracy: 0.7855

Training time: 2489.17 seconds





Test predictions:

Text: Scientists cure cancer, world peace achieved

Dense Network: 0.3318 (Negative)

LSTM Network: 0.0367 (Negative)

WARNING:tensorflow:5 out of the last 50 calls to <function TensorFlowTrainer.make\_predict\_function.<locals>.one\_step\_on\_dat

Bidirectional LSTM Network: 0.3236 (Negative)

WARNING:tensorflow:6 out of the last 51 calls to <function TensorFlowTrainer.make\_predict\_function.<locals>.one\_step\_on\_dat

Double Bidirectional LSTM Network: 0.3377 (Negative)

Text: Local man wins lottery, plans to spend it all on cat food

Dense Network: 0.9577 (Positive)

LSTM Network: 0.9713 (Positive)

Bidirectional LSTM Network: 0.9267 (Positive)

Double Bidirectional LSTM Network: 0.9526 (Positive)

Text: Breaking: Water is wet, scientists confirm

Dense Network: 0.7731 (Positive)

LSTM Network: 0.8613 (Positive)

Bidirectional LSTM Network: 0.8451 (Positive)

Double Bidirectional LSTM Network: 0.9532 (Positive)

Model Performance Summary:

	final_train_acc	final_val_acc	\
Dense Network	0.8570	0.7935	
LSTM Network	0.9440	0.7851	
Bidirectional LSTM Network	0.8531	0.7862	
Double Bidirectional LSTM Network	0.9250	0.7855	

	best_val_acc	final_train_loss	\
Dense Network	0.7945	0.3804	
LSTM Network	0.7855	0.1833	
Bidirectional LSTM Network	0.7862	0.3683	
Double Bidirectional LSTM Network	0.7917	0.2232	

	final_val_loss	training_time
Dense Network	0.4564	122.1702
LSTM Network	0.6601	2420.8608
Bidirectional LSTM Network	0.4608	1172.9664
Double Bidirectional LSTM Network	0.5459	2489.1669

## 2. The IMDB dataset

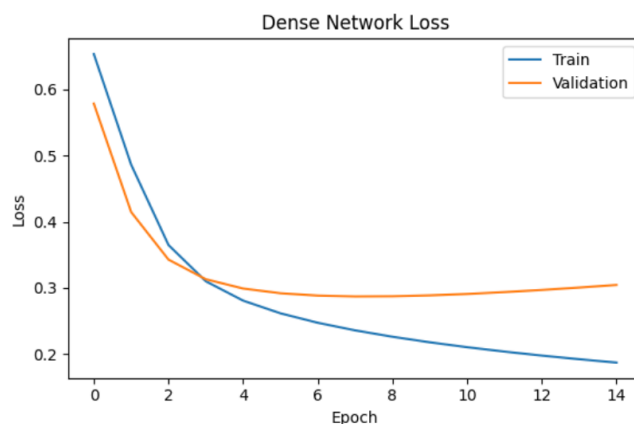
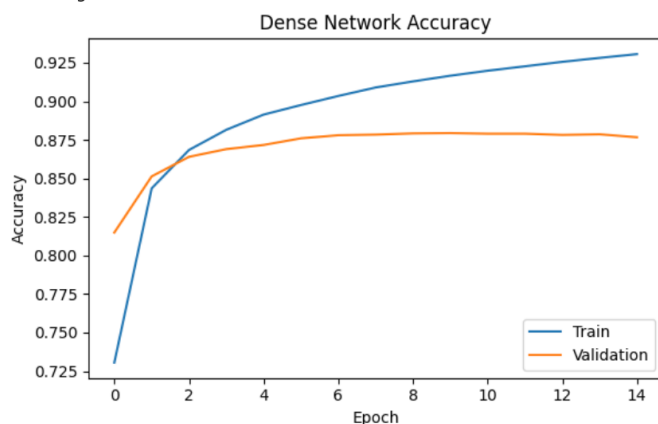
Dense:

```

Epoch 12/15 1250/1250 8s 7ms/step - accuracy: 0.9200 - loss: 0.2000 - val_accuracy: 0.8750 - val_loss: 0.2910
Epoch 13/15 1250/1250 9s 5ms/step - accuracy: 0.9217 - loss: 0.2017 - val_accuracy: 0.8790 - val_loss: 0.2938
Epoch 14/15 1250/1250 8s 7ms/step - accuracy: 0.9244 - loss: 0.1956 - val_accuracy: 0.8782 - val_loss: 0.2970
Epoch 15/15 1250/1250 8s 7ms/step - accuracy: 0.9275 - loss: 0.1900 - val_accuracy: 0.8786 - val_loss: 0.3006
Epoch 15/15 1250/1250 7s 5ms/step - accuracy: 0.9297 - loss: 0.1849 - val_accuracy: 0.8767 - val_loss: 0.3046

```

Dense Network Validation Accuracy: 0.8767  
 Training time: 136.06 seconds



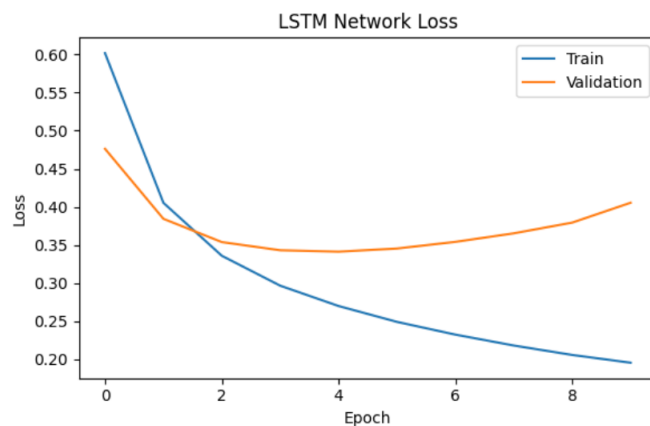
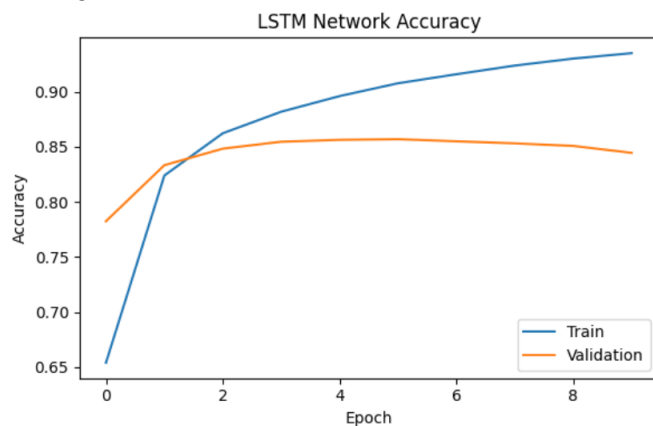
## LSTM:

```

Epoch 8/10 1250/1250 250s 166ms/step - accuracy: 0.9242 - loss: 0.2186 - val_accuracy: 0.8532 - val_loss: 0.3652
Epoch 9/10 1250/1250 202s 161ms/step - accuracy: 0.9315 - loss: 0.2056 - val_accuracy: 0.8509 - val_loss: 0.3793
Epoch 10/10 1250/1250 201s 161ms/step - accuracy: 0.9366 - loss: 0.1943 - val_accuracy: 0.8446 - val_loss: 0.4054

```

LSTM Network Validation Accuracy: 0.8446  
 Training time: 2080.61 seconds



Training Bidirectional LSTM Network...

Embedding dimension: 64

Using learning rate: 1e-05 for Bidirectional LSTM Network

Training for 10 epochs

Epoch 1/10

1250/1250 ————— 141s 110ms/step - accuracy: 0.5048 - loss: 0.6930 - val\_accuracy: 0.5121 - val\_loss: 0.6922

Epoch 2/10

1250/1250 ————— 140s 108ms/step - accuracy: 0.5401 - loss: 0.6857 - val\_accuracy: 0.7355 - val\_loss: 0.5492

Epoch 3/10

1250/1250 ————— 141s 108ms/step - accuracy: 0.7921 - loss: 0.4917 - val\_accuracy: 0.8318 - val\_loss: 0.4162

Epoch 4/10

1250/1250 ————— 142s 108ms/step - accuracy: 0.8408 - loss: 0.3910 - val\_accuracy: 0.8590 - val\_loss: 0.3739

Epoch 5/10

1250/1250 ————— 140s 106ms/step - accuracy: 0.8686 - loss: 0.3314 - val\_accuracy: 0.8702 - val\_loss: 0.3231

Epoch 6/10

1250/1250 ————— 144s 108ms/step - accuracy: 0.8935 - loss: 0.2745 - val\_accuracy: 0.8749 - val\_loss: 0.3134

Epoch 7/10

1250/1250 ————— 136s 109ms/step - accuracy: 0.9066 - loss: 0.2479 - val\_accuracy: 0.8760 - val\_loss: 0.3128

Epoch 8/10

1250/1250 ————— 138s 106ms/step - accuracy: 0.9170 - loss: 0.2280 - val\_accuracy: 0.8759 - val\_loss: 0.3168

Epoch 9/10

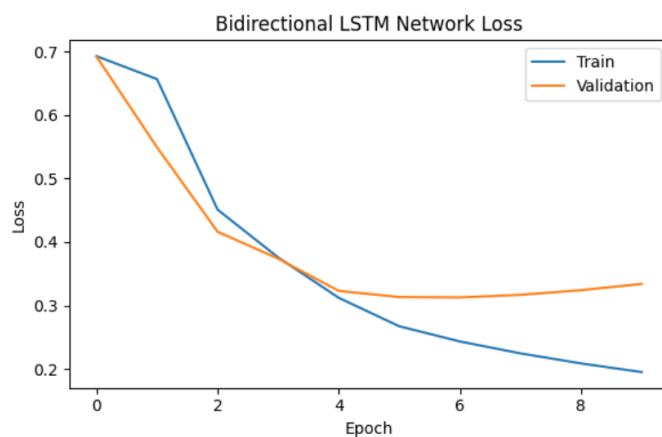
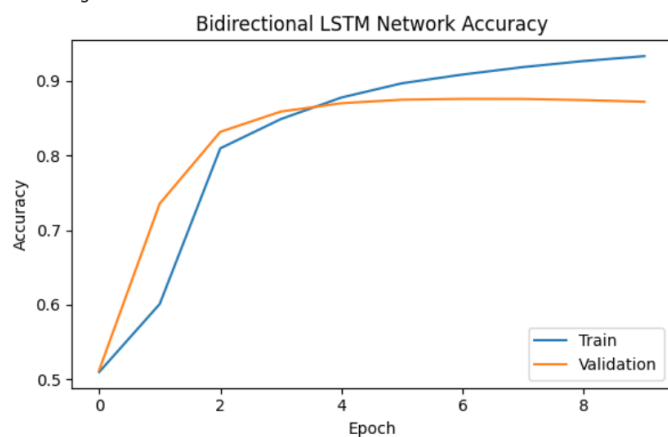
1250/1250 ————— 144s 108ms/step - accuracy: 0.9259 - loss: 0.2116 - val\_accuracy: 0.8744 - val\_loss: 0.3242

Epoch 10/10

1250/1250 ————— 142s 108ms/step - accuracy: 0.9328 - loss: 0.1977 - val\_accuracy: 0.8721 - val\_loss: 0.3340

Bidirectional LSTM Network Validation Accuracy: 0.8721

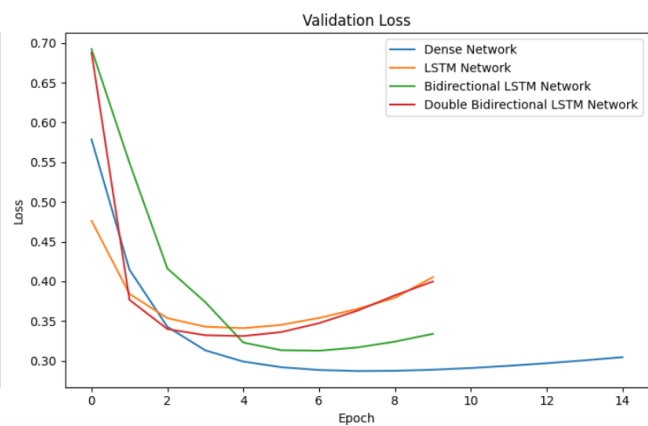
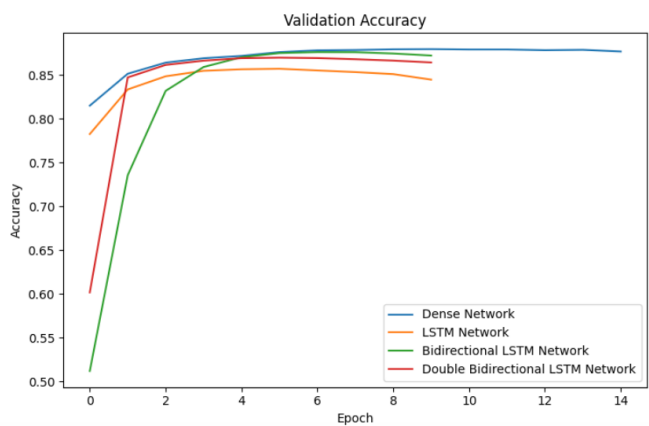
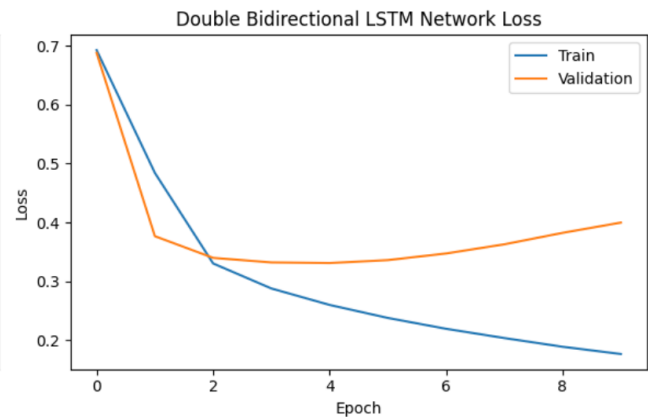
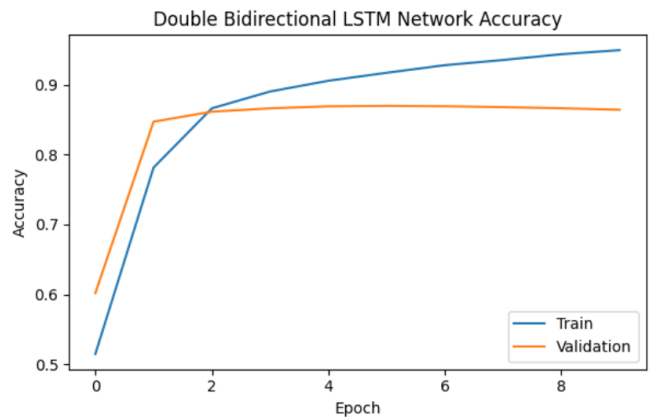
Training time: 1416.51 seconds





1250/1250 — 323s 228ms/step — accuracy: 0.9353 — loss: 0.2017 — val\_accuracy: 0.8679 — val\_loss: 0.3630  
Epoch 9/10  
1250/1250 — 286s 229ms/step — accuracy: 0.9443 — loss: 0.1861 — val\_accuracy: 0.8663 — val\_loss: 0.3824  
Epoch 10/10  
1250/1250 — 285s 228ms/step — accuracy: 0.9501 — loss: 0.1723 — val\_accuracy: 0.8642 — val\_loss: 0.3998

Double Bidirectional LSTM Network Validation Accuracy: 0.8642  
Training time: 3070.10 seconds



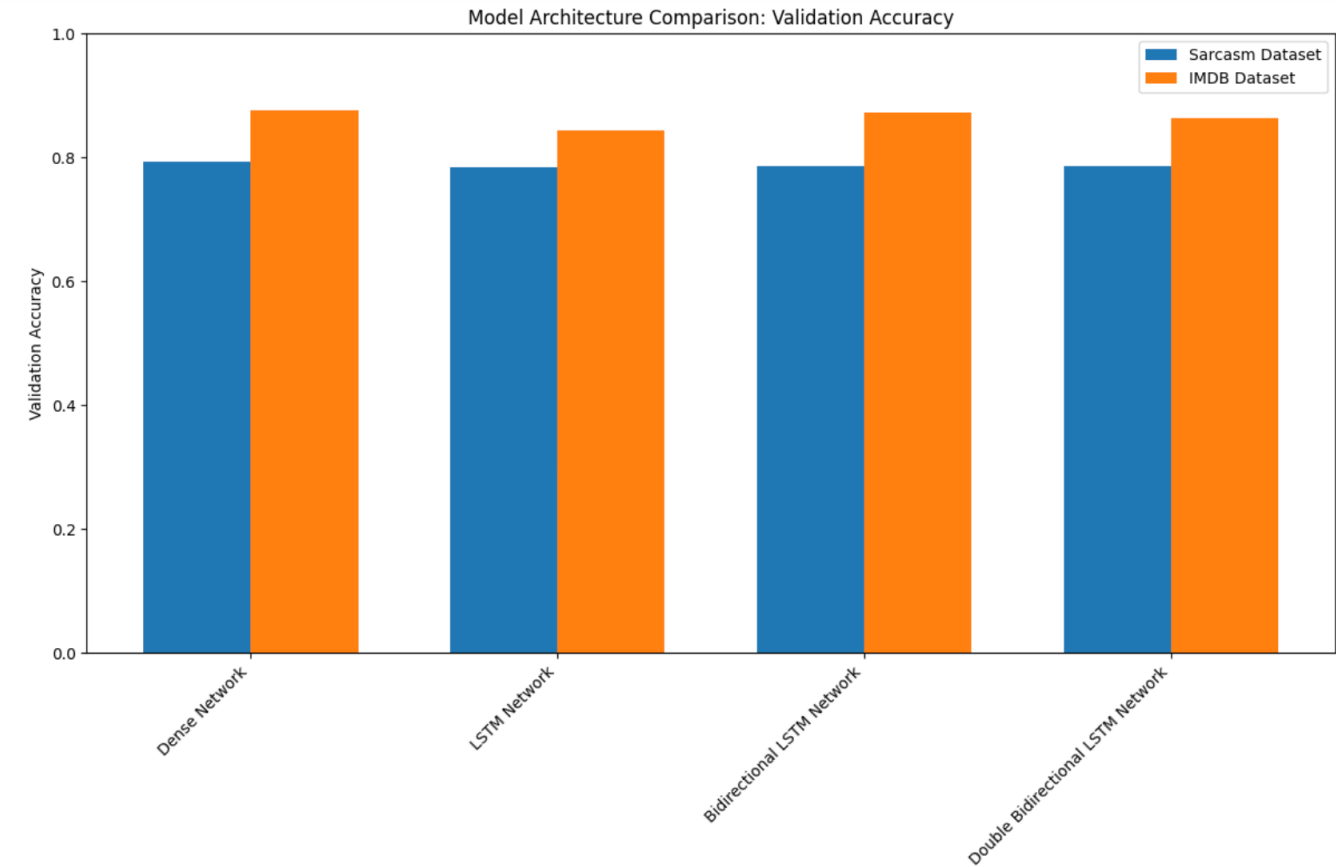
Test predictions:

Text: this film was absolutely terrible the worst acting i have seen  
Dense Network: 0.0741 (Negative)  
LSTM Network: 0.0179 (Negative)  
Bidirectional LSTM Network: 0.0800 (Negative)  
Double Bidirectional LSTM Network: 0.0227 (Negative)

Text: a masterpiece of cinema with outstanding performances and direction  
Dense Network: 0.9119 (Positive)  
LSTM Network: 0.9879 (Positive)  
Bidirectional LSTM Network: 0.9244 (Positive)  
Double Bidirectional LSTM Network: 0.9609 (Positive)

Text: neither good nor bad just a mediocre film that passes the time  
Dense Network: 0.2011 (Negative)  
LSTM Network: 0.0475 (Negative)  
Bidirectional LSTM Network: 0.0930 (Negative)  
Double Bidirectional LSTM Network: 0.0606 (Negative)

Model Performance Summary:			
	final_train_acc	final_val_acc	\
Dense Network	0.9306	0.8767	
LSTM Network	0.9352	0.8446	
Bidirectional LSTM Network	0.9333	0.8721	
Double Bidirectional LSTM Network	0.9495	0.8642	
	best_val_acc	final_train_loss	\
Dense Network	0.8794	0.1872	
LSTM Network	0.8570	0.1953	
Bidirectional LSTM Network	0.8760	0.1952	
Double Bidirectional LSTM Network	0.8697	0.1769	
	final_val_loss	training_time	
Dense Network	0.3046	136.0565	
LSTM Network	0.4054	2080.6131	
Bidirectional LSTM Network	0.3340	1416.5087	
Double Bidirectional LSTM Network	0.3998	3070.1048	



Comparison complete! Review the results and charts above.

## Reflections:

### 1. Model Complexity vs. Performance

The Dense Network outperformed more complex architectures on both datasets. What might explain this counterintuitive result, and what implications does this have for the principle that "more complex models yield better results"?

**Ans:** Although Dense is "simpler" than the more complexed LSTM model, possibly the dataset of the IMDB doesn't have time-series characteristics, and using Dense would be enough and more efficient. So it is not true that complex models always yield better result. There's a trade-off between complexity and efficiency.

### 2. Training-Validation Gap Analysis

The LSTM Network on the sarcasm dataset showed a large gap between training accuracy and validation accuracy. What does this pattern suggest, and what techniques could you implement to address this issue?

**Ans:** It shows the over-fitting problem. To mitigate the over-fitting problem, we can add the Dropout layer, or we can set the early stopping.

### 3. Task Difficulty Comparison

All models performed better on the IMDB dataset than on the sarcasm dataset. What characteristics of sarcasm might make it inherently more difficult to detect compared to sentiment analysis?

**Ans:** I think a possible reason is that the reviews of IMDB have stronger or more intense sentiments, while the sentences on the sarcasm dataset have more ambiguous emotions. So it makes the sarcasm dataset harder to detect.

### 4. Computational Efficiency Analysis

The training times varied dramatically between architectures. Considering the performance results, analyze the trade-offs between model complexity, computational resources, and accuracy. When might you choose a simpler model despite having resources for more complex ones?

**Ans:** A more complex model may require larger computational resources (and a lot of time to train the models), like the LSTMs in this example, and when the dataset itself is large and complex, a more sophisticated model may bring higher accuracy. However, sometimes when the data itself is simple, and a simple model can provide good accuracy. In this case, we don't need to use a very complex model, even if we have enough resources.

### 5. Architecture Enhancement Proposal

Based on these results, propose a modified architecture or approach that might improve performance on the sarcasm detection task while maintaining reasonable com-

putational efficiency. Justify your design choices.

**Ans:** Perhaps I would include the Attention model, which is introduced in the following lectures. This design may be better than having multiple layers of LSTM.

### **Part3: Dense vs LSTM on longer documents**

Sorry that I haven't done this part yet!